

# Markov Decision Processes with Dynamic Transition Probabilities: An Analysis of Shooting Strategies in Basketball

A Simplified Walkthrough in R

*Nathan Sandholtz and Luke Bornn*

*3/2/2020*

## Data

The folder `./data` contains one game of optical player tracking data from the 2012-2013 NBA regular season filtered to observations with tagged ball-events including dribbles, passes, turnovers, and shots. The game was between the Miami Heat and the Brooklyn Nets. Additionally, we filtered the data to plays in which no fouls occurred. For this simplified walkthrough, we have categorized players into three position groups: Guards (G), Forwards (F), and Centers (C).

```
# Load data
dat = read.csv("./data/2013_11_01_MIA_BRK_formatted.csv")
head(dat)
```

##	game	time	quarter	game_clock	shot_clock	entity	team	x
## 1	2013110117	98398	1	675.56	24.00	172537	BRK	46.88853
## 2	2013110117	99198	1	674.76	23.20	172537	BRK	42.26108
## 3	2013110117	99358	1	674.60	23.04	172537	BRK	41.44996
## 4	2013110117	100278	1	673.68	22.12	172537	BRK	36.87506
## 5	2013110117	101318	1	672.64	21.08	172537	BRK	32.91451
## 6	2013110117	101998	1	671.96	20.40	172537	BRK	32.55346

##	y	event_id	ndd	change_poss	play	terminal_event	location_id
## 1	11.02848	21	10.408324	1	2	FALSE	heave
## 2	11.64887	21	8.325748	0	2	FALSE	heave
## 3	11.88701	21	7.790360	0	2	FALSE	heave
## 4	13.75815	21	4.816156	0	2	FALSE	heave
## 5	16.13047	21	1.690733	0	2	FALSE	heave
## 6	17.08186	21	3.705388	0	2	FALSE	heave

##	def_pres	points	time_lapse	firstname	lastname	position_simple	home_away
## 1	open	0	-0.80	Deron	Williams	G	h
## 2	open	0	-0.16	Deron	Williams	G	h
## 3	open	0	-0.92	Deron	Williams	G	h
## 4	contested	0	-1.04	Deron	Williams	G	h
## 5	contested	0	-0.68	Deron	Williams	G	h
## 6	contested	0	-0.24	Deron	Williams	G	h

We also source some utility functions that will be used in the walkthrough. These functions include Algorithm 1 from the paper and functions used to get the initial states and shot clock times for each team's set of plays and the empirical distribution of time lapses between on-ball events.

```
# Source utils
source("./code/simulation_utils.R")
```

```
##
## Attaching package: 'dplyr'
```

```
## The following objects are masked from 'package:stats':
##
##   filter, lag
## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union
```

## Models

In order to simulate plays, we require fits from the models of the latent MDP componets. Specifically, the player-specific parameters from the estimated shot policy model, state transition model, and reward model. These models are formally defined in Section 3 of the paper. We fit the models using Stan. The Stan model scripts are included in the folder `./code/stan_models` and the code to fit them are contained in files `./code/policy_fit.R`, `./code/transition_probability_fit.R`, and `./code/reward_fit.R` respectively. As these models take a considerable amount of time to fit, we have included 300 posterior draws from the player-specific parameters for each model.

```
# Load posterior draws
n_draws = 300
lambda_MIA_draws = readRDS("./model_output/lambda_MIA_draws.rds")
lambda_BRK_draws = readRDS("./model_output/lambda_BRK_draws.rds")
mu_draws = readRDS("./model_output/mu_draws.rds")
theta_draws = readRDS("./model_output/theta_draws.rds")
xi_draws = readRDS("./model_output/xi_draws.rds")
```

## Simulating plays

Before we can simulate plays, we require a few more in inputs. As noted in Section 4 of the paper, we require the starting states of all plays and the corresponding shot clock times at the start of each play. We also need the empirical distribution of time-lapses between events in order to take time off of the shot clock at each step of the MDP.

```
# Get initial states and shot clock times for each team.
MIA_initial_states <- get_initial_states(dat, "MIA")
BRK_initial_states <- get_initial_states(dat, "BRK")

# Get empirical shot clock distribution
shot_clock_dist <- get_sc_dist(dat = dat, num_intervals = 3)
```

We can now simulate each team's plays in this game for a chosen number of simulations. We'll simulate the game 100 times.

```
n_sim = 100

# MIAMI SIMULATIONS
MIA_points = NA
for(iter in 1:n_sim){
  for(play in 1:nrow(MIA_initial_states)) {
    if (play == 1) {
      game_moments_MIA = algorithm_1(
        s_0 = MIA_initial_states[play, "state"],
        c_0 = MIA_initial_states[play, "shot_clock"],
        theta_draws = theta_draws,
```

```

        mu_draws = mu_draws,
        xi_draws = xi_draws,
        lambda_draws = lambda_MIA_draws,
        L_dist = shot_clock_dist,
        num_mcmc = n_draws
    )
} else {
    game_moments_MIA = rbind(
        game_moments_MIA,
        algorithm_1(
            s_0 = MIA_initial_states[play, "state"],
            c_0 = MIA_initial_states[play, "shot_clock"],
            theta_draws = theta_draws,
            mu_draws = mu_draws,
            xi_draws = xi_draws,
            lambda_draws = lambda_MIA_draws,
            L_dist = shot_clock_dist,
            num_mcmc = n_draws
        )
    )
}
MIA_points[iter] = sum(game_moments_MIA$reward)
}

```

#### *# BROOKLYN SIMULATIONS*

```

BRK_points = NA
for(iter in 1:n_sim){
    for(play in 1:nrow(BRK_initial_states)) {
        if (play == 1) {
            game_moments_BRK = algorithm_1(
                s_0 = BRK_initial_states[play, "state"],
                c_0 = BRK_initial_states[play, "shot_clock"],
                theta_draws = theta_draws,
                mu_draws = mu_draws,
                xi_draws = xi_draws,
                lambda_draws = lambda_BRK_draws,
                L_dist = shot_clock_dist,
                num_mcmc = n_draws
            )
        } else {
            game_moments_BRK = rbind(
                game_moments_BRK,
                algorithm_1(
                    s_0 = BRK_initial_states[play, "state"],
                    c_0 = BRK_initial_states[play, "shot_clock"],
                    theta_draws = theta_draws,
                    mu_draws = mu_draws,
                    xi_draws = xi_draws,
                    lambda_draws = lambda_BRK_draws,
                    L_dist = shot_clock_dist,
                    num_mcmc = n_draws
                )
            )
        }
    }
}

```

```

    )
  }
}
BRK_points[iter] = sum(game_moments_BRK$reward)
}

```

We can plot density estimates of our simulations and compare these to the empirical total points from these plays in the data. Dotted vertical lines represent each team's observed total points from the filtered plays from this game.

```

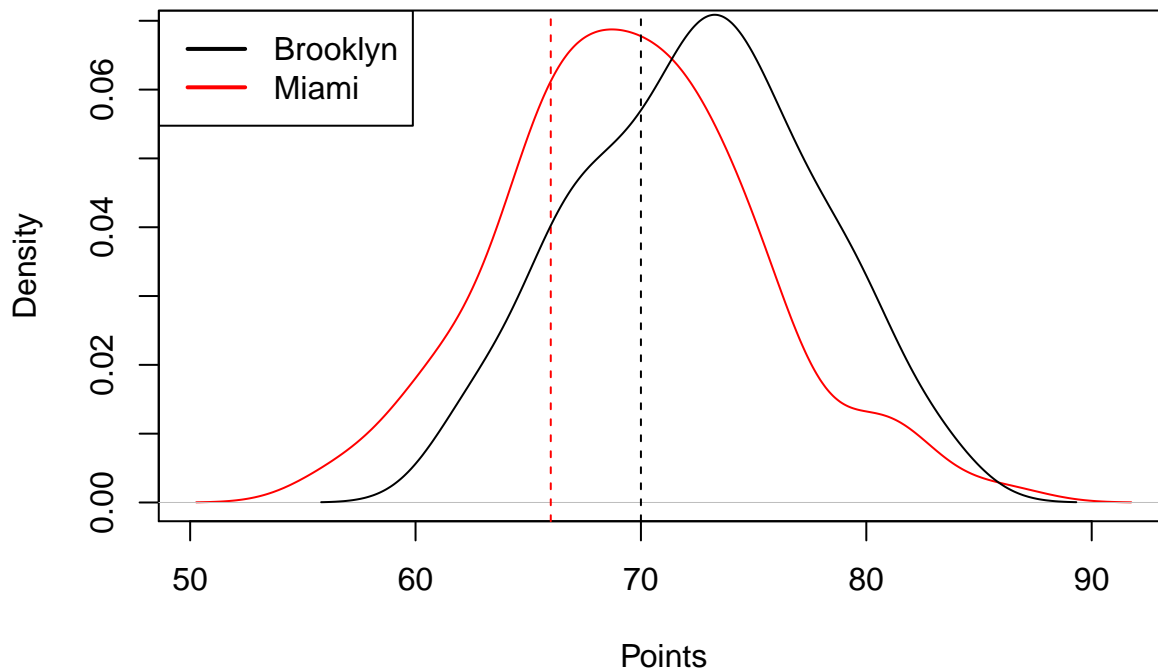
# Compare simulations to empirical
plot(density(MIA_points), col = "red",
     main = "Simulations: MIA vs BRK",
     xlab = "Points")
dat %>% filter(team == "MIA") %>% with(abline(v = sum(points),
                                             col = "red",
                                             lty = 2))

lines(density(BRK_points))
dat %>% filter(team == "BRK") %>% with(abline(v = sum(points),
                                             lty = 2))

legend("topleft", c("Brooklyn", "Miami"),
      col = c("black", "red"),
      lwd = 2, bg = NA)

```

## Simulations: MIA vs BRK



## Altering Policies

We now will explore an alteration to Miami's shot policy. We will decrease each player's midrange shot policy by 20% (except late in shot clock) and increase each player's three point policy by 20% regardless of time on

clock. The function `alter_theta` contained in the simulation utilities script alters the posterior draws of `theta` according to our desired changes.

```
# POLICY ALTERATION
# Decrease midrange shot policy by 20% (except late in shot clock) and
# increase three point policy by 20% (regardless of time on clock)

# Identify MIA players
MIA_players = dat %>%
  filter(team == "MIA") %>%
  distinct(entity) %>%
  pull(entity)

# Identify states to alter
# 1) ALL Midrange shots
to_alter_1 = c(paste(MIA_players, "long2_contested", sep = "_"),
              paste(MIA_players, "long2_open", sep = "_"))
# 2) ALL three point shots
to_alter_2 = c(paste(MIA_players, "three_contested", sep = "_"),
              paste(MIA_players, "three_open", sep = "_"))

policy_change <- list(list(who_where = to_alter_1,
                          when = 2:3,
                          how_much = .8),
                     list(who_where = to_alter_2,
                          when = 1:3,
                          how_much = 1.2)
                     )

# Alter the posterior draws of theta
altered_theta_draws = alter_theta(theta_draws,
                                  altered_policy_rules = policy_change)

# MIAMI ALTERED SIMULATIONS
MIA_points_alt = NA
for(iter in 1:n_sim){
  for(play in 1:nrow(MIA_initial_states)) {
    if (play == 1) {
      game_moments_MIA = algorithm_1(
        s_0 = MIA_initial_states[play, "state"],
        c_0 = MIA_initial_states[play, "shot_clock"],
        theta_draws = altered_theta_draws,
        mu_draws = mu_draws,
        xi_draws = xi_draws,
        lambda_draws = lambda_MIA_draws,
        L_dist = shot_clock_dist,
        num_mcmc = n_draws
      )
    } else {
      game_moments_MIA = rbind(
        game_moments_MIA,
        algorithm_1(
          s_0 = MIA_initial_states[play, "state"],
```


```

    c_0 = MIA_initial_states[play, "shot_clock"],
    theta_draws = altered_theta_draws,
    mu_draws = mu_draws,
    xi_draws = xi_draws,
    lambda_draws = lambda_MIA_draws,
    L_dist = shot_clock_dist,
    num_mcmc = n_draws
  )
}
}
MIA_points_alt[iter] = sum(game_moments_MIA$reward)
}

plot(density(MIA_points), col = "red",
     main = "Simulations: MIA vs BRK",
     xlab = "Points")
lines(density(BRK_points))
lines(density(MIA_points_alt), col = "blue")
legend("topleft", c("Brooklyn", "Miami", "Miami-alt"),
     col = c("black", "red", "blue"),
     lwd = 2)

```

### Simulations: MIA vs BRK



Miami's projected distribution of possible scores increases.